

IERG4130

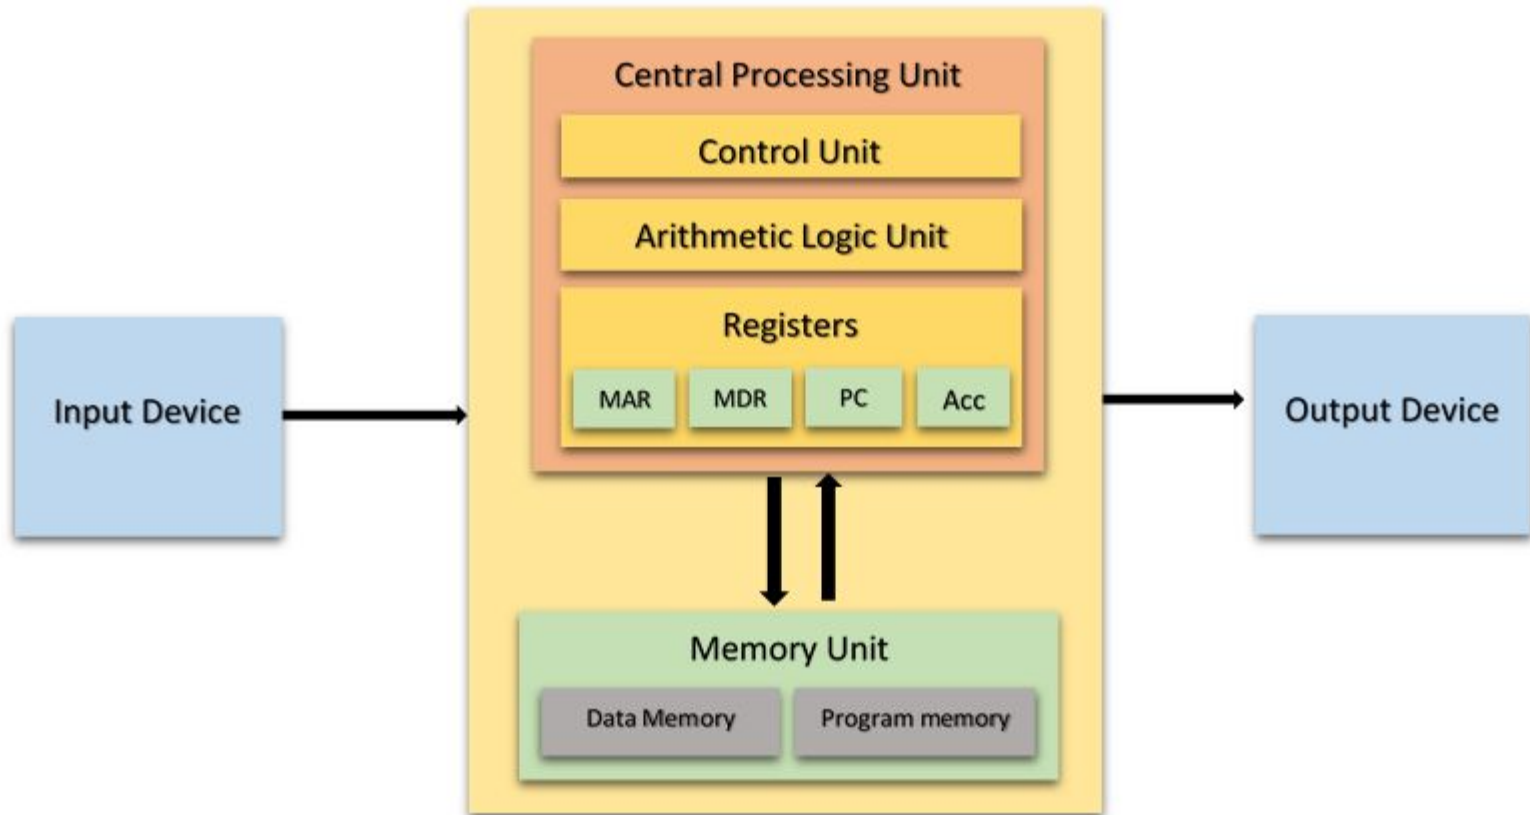
Buffer Overflow

By Ke ZHANG
20/9/2022

Outline

- Background
 - Memory Layout
 - Stack Layout
 - What happens when a function is called
 - Calling Conventions
- Buffer Overflow
 - How to launch
 - How to mitigate
- Example

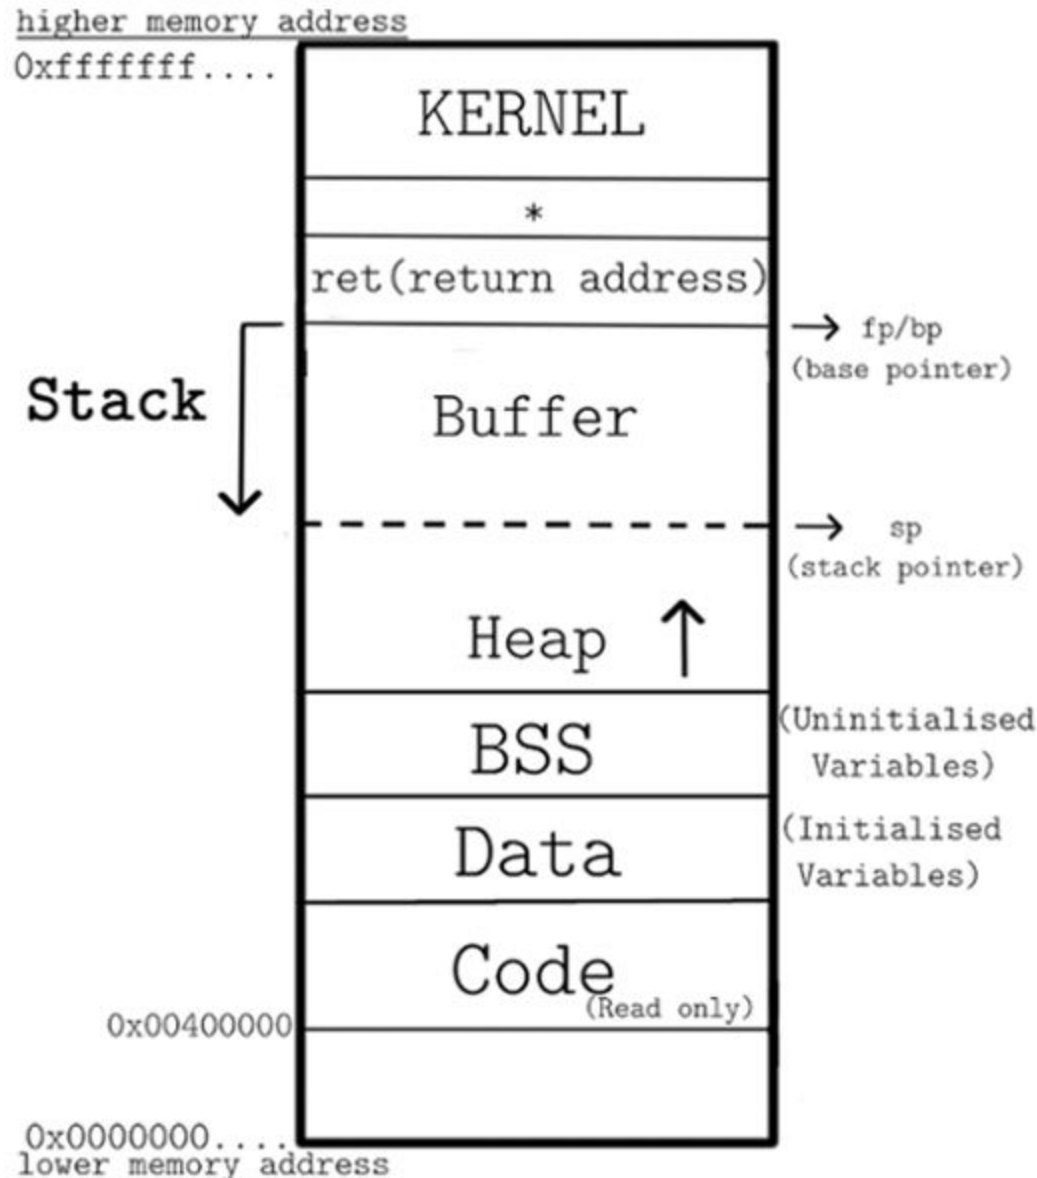
Assumption: Von-Neumann Architecture



Memory Layout


- Program and Process?
 - Program: static; a set of instructions; stored as a file
 - Process: dynamic; when a program runs, it becomes a ~; holds resources
 - A process exists in **memory**
- Memory
 - Memory address space
 - User space vs Kernel space
 - Virtual memory space vs Physical memory space

Memory Organization



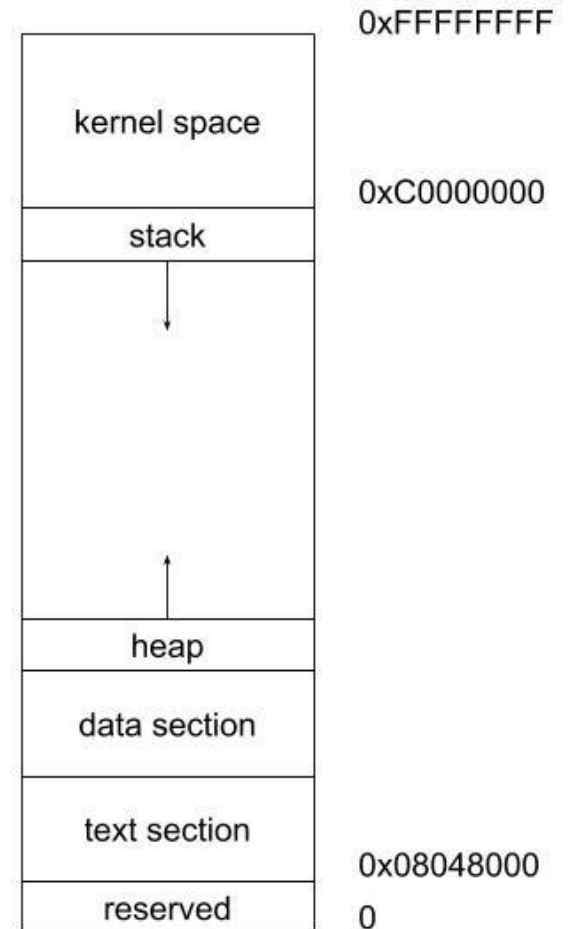
Virtual Memory

- Each process has its own **virtual** memory
 - virtual memory maps to physical memory
 - $\text{size}(\text{virtual memory}) \gg \text{size}(\text{physical memory})$
 - handle by MMU (Memory Management Unit)
 - Page and Page frame
 - Page table
 - Why we need virtual memory?
 - Abstraction for unified layout
 - Ignore real physical address
 - Isolation
 - Sharing



Memory Layout

- (Here we only talk about Virtual Memory)
- Stack:
 - 1. data structure: LIFO
 - 2. memory area:
 - grow from high addr to low addr
 - destroy from top to bottom (base)
- For a **function** (procedure)
 - Stack Frame (Activity Record)
 - What does it contain?




Stack Frame

what is it all about??



- (We have come from process-level to function-level)
- You may have a lot of questions:
 - what is EBP and ESP?
 - what does 'old' means in 'old EBP'?
 - return address? return where?
 -



Stack Frame

- EBP: Base Pointer, pointing to the stack bottom (high address)
 - fixed
 - $EBP + \text{offset}$: to locate variables
- ESP: Stack Pointer, pointing to the stack top (low address)
 - shifted when POP & PUSH
 - $ESP - EBP$: the space of a stack
- Note that
 - both EBP and ESP are in CPU.
 - Old EBP and saved registers are values in main memory.

EBP and ESP

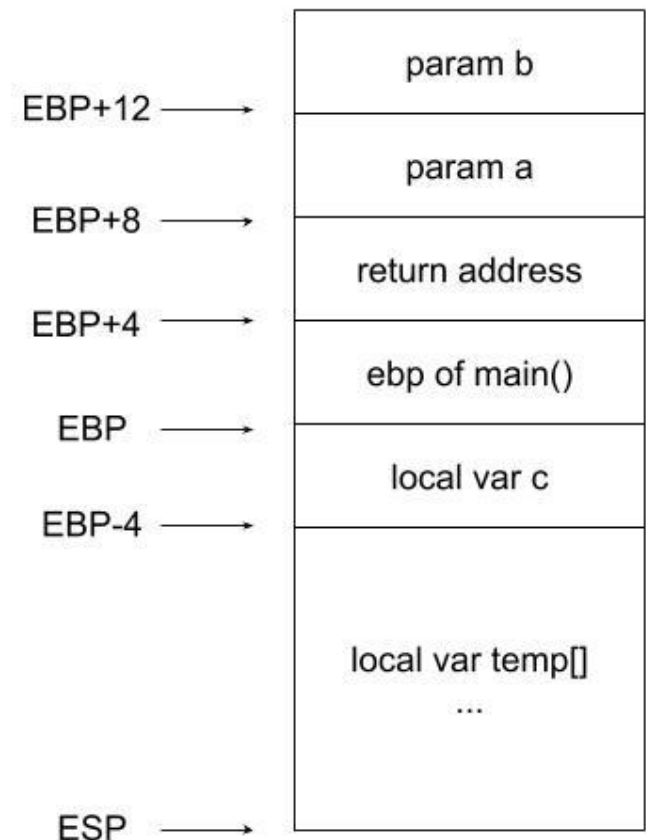
- Use the sample code in lecture notes

```
- #include <stdio.h>
int add(int a, int b) {
    int c, temp[512];
    c = a + b;
    return c;    }

int main(int argc, char * argv) {
    int one, two, three;
    one = 1; two = 2;
    three = add(one, two);
    return three;
}
```

EBP and ESP

- For ***add*** function called by main()
- (assume we do not store registers in stack)
 - for simplicity
 - also not consider stack alignment here
- (assume we push parameters from right to left in source code)
 - another question → why? how?
 - explain in later slides
- [EBP + offset] to locate element inside current frame.

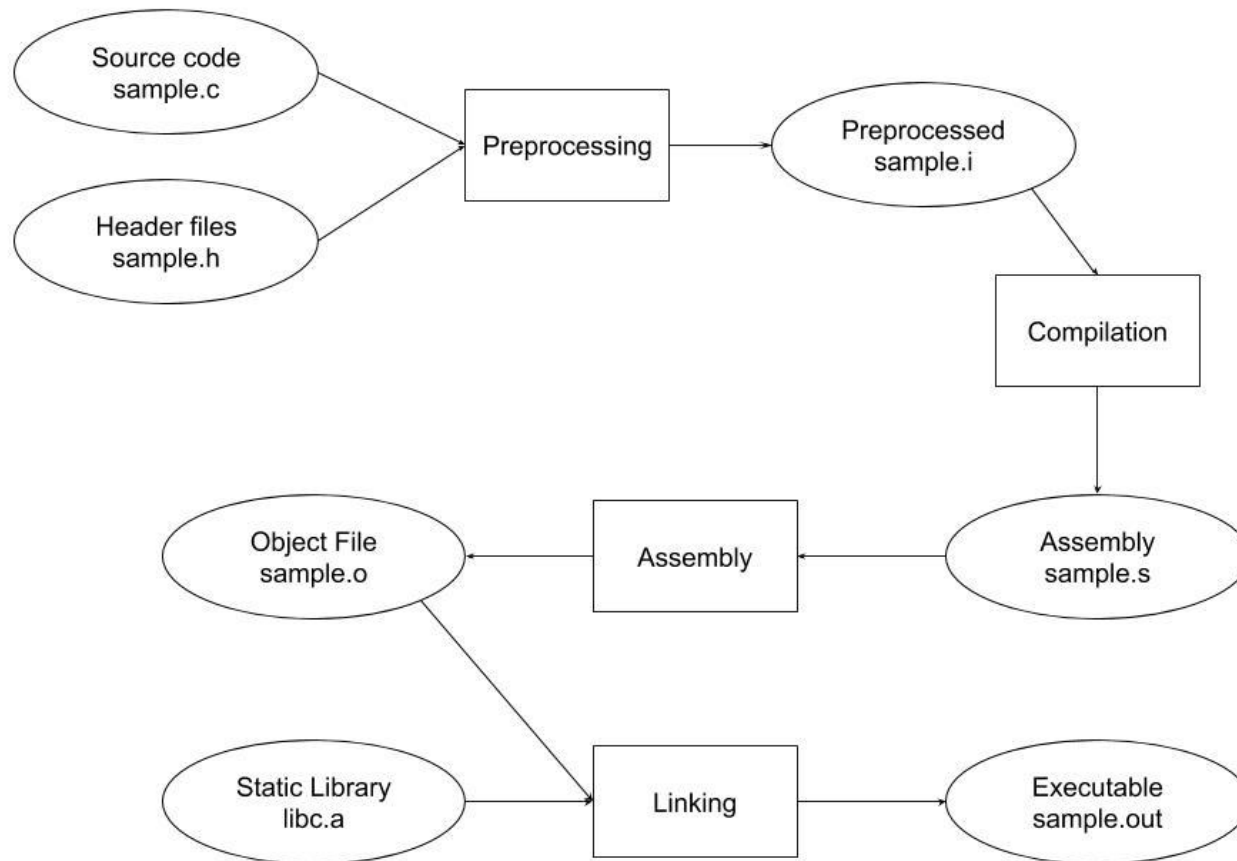


Stack Frame

- Old EBP & Return Address
 - related to function calls
- Caller & Callee
 - calling function → caller; called function → callee
 - Stack Frame (also known as Activity Record) is function-level
 - Each function has its own stack frame
- When funcA calls funcB, funcB calls funcC ...
 - A chain
 - When funcC returns, we need trace back to funcB, and so on
- old EBP: to recover caller's stack frame
- return address: to continue the exec of caller
 - the next instruction to exec after the callee's return

Assembly basics

- Think about how you get an executable from C source code.



Assembly basics

- Registers (in CPU)


- General registers
 - Data registers
 - Pointer registers
 - Index registers
- Control registers
- Segment registers

- Data registers

- EAX, EBX, ECX, EDX

- Pointer registers

- EIP (Instruction Pointer): stores the offset address of the next instruction to be executed
- ESP (Stack Pointer) - the top element of stack
- EBP (Base Pointer) - the stack base



Assembly basics

Note the following illustrations are simplified to let you familiar with their functions, for a detailed reference please check official manuals.

- Syntax: **Intel** v.s. **AT&T**
 - **MOV** EAX, EBX; transfer data in EBX to EAX
 - **ADD** EAX, EBX; add data in EAX and data in EBX and store the result to EAX
 - **PUSH** EAX;
 - from register (CPU) to memory
 - equal to $ESP=ESP-4$; **MOV** [ESP], EAX
 - **POP** EAX;
 - from memory to register (CPU).
 - equal to **MOV** EAX, [ESP]; $ESP=ESP+4$
 - **CALL** func;
 - equal to **PUSH** EIP; **JMP** func
 - **RET**; equal to **POP** EIP

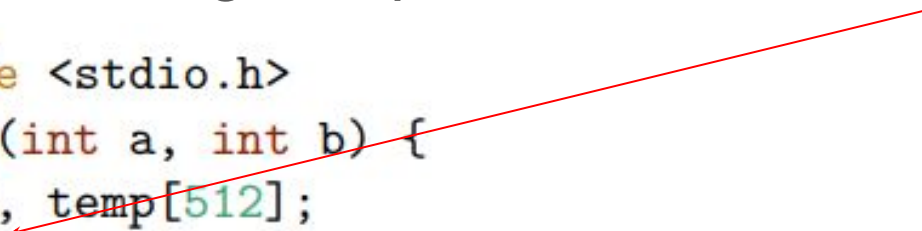
Sample

- Use the following sample code

```
- #include <stdio.h>
int add(int a, int b) {
    int c, temp[512];
    c = a + b;
    return c;    }

int main(int argc, char * argv) {
    int one, two, three;
    one = 1; two = 2;
    three = add(one, two);
    return three;
}
```

add an instruction:
temp[0] = 10;



What happens in **caller**: main()

```
main:
.LFB1:
    .loc 1 9 0
    .cfi_startproc
    push    ebp
    .cfi_def_cfa_offset 8
    .cfi_offset 5, -8
    mov     ebp, esp
    .cfi_def_cfa_register 5
    sub     esp, 16
    .loc 1 11 0
    mov     DWORD PTR [ebp-4], 1
    mov     DWORD PTR [ebp-8], 2
    .loc 1 12 0
    push    DWORD PTR [ebp-8]
    push    DWORD PTR [ebp-4]
    call    add
    add     esp, 8
    mov     DWORD PTR [ebp-12], eax
    .loc 1 13 0
    mov     eax, DWORD PTR [ebp-12]
    .loc 1 14 0
    leave
    .cfi_restore 5
    .cfi_def_cfa 4, 4
    ret
    .cfi_endproc
```

Every function starts with
push ebp;
mov ebp, esp

allocate space for its stack

for main(), 1 and 2 are **local variables**

push 1 and 2 to the stack
push 1=> ESP=ESP-4; [ESP]=1

remove all pushed parameters,
so that we can restore those
saved registers

What happens in **callee**: add()

```
add:
.LFB0:
.file 1 "sample.c"
.loc 1 2 0
.cfi_startproc
push    ebp
.cfi_def_cfa_offset 8
.cfi_offset 5, -8
mov     ebp, esp
.cfi_def_cfa_register 5
sub     esp, 2064
.loc 1 4 0
mov     DWORD PTR [ebp-2052], 10
.loc 1 5 0
mov     edx, DWORD PTR [ebp+8]
mov     eax, DWORD PTR [ebp+12]
add     eax, edx
mov     DWORD PTR [ebp-4], eax
.loc 1 6 0
mov     eax, DWORD PTR [ebp-4]
.loc 1 7 0
leave
.cfi_restore 5
.cfi_def_cfa 4, 4
ret
.cfi_endproc
```

Every function starts with *push ebp; mov ebp, esp*

I add "temp[0] = 10;" in source code

parameter a

parameter b

local variable c

put return value in eax

restore the old ebp
equals to:

mov esp,ebp
pop ebp

ret equals to
pop eip

Calling Conventions

- Caller and callee agree on a convention
 - the order of pushing parameters (left to right? right to left?)
 - how to balance the stack
 - callee or caller?
 - name mangling
 - for linking stage
- `__cdecl`
 - C default calling convention → from right to left; caller to balance
- `__stdcall`
 - win32 API → from right to left; callee to balance

cdecl and stdcall

```
_cdecl int MyFunction1(int a, int b)
{
    return a + b;
}
```

and the following function call:

```
x = MyFunction1(2, 3);
```

```
_MyFunction1:
push ebp
mov ebp, esp
mov eax, [ebp + 8]
mov edx, [ebp + 12]
add eax, edx
pop ebp
ret
```

and

```
push 3
push 2
call MyFunction1
add esp, 8
```

```
_stdcall int MyFunction2(int a, int b)
{
    return a + b;
}
```

and the calling instruction:

```
x = MyFunction2(2, 3);
```

```
:_MyFunction2@8
push ebp
mov ebp, esp
mov eax, [ebp + 8]
mov edx, [ebp + 12]
add eax, edx
pop ebp
ret 8
```

and

```
push 3
push 2
call _MyFunction2@8
```

Buffer Overflow

- Why?
 - for continuous memory writing
 - no boundary check for data on stack
- Severe?
 - Overwrite data on stack
 - hijack the control flow → arbitrary (malicious) code execution
- How to launch?
 - Step 1. Find those standard string function calls without a compulsory boundary check
 - A lot of C standard functions
 - strcmp(), strcpy(), ...

How to launch buffer overflow attack

- step **2**. Write a shellcode
 - Payload
 - Why this name?
- Let's see an example
 - gcc -fno-stack-protector -ggdb bufferoverflow.c -o bufferoverflow
 - gdb ./bufferoverflow
 - str='A'*24

[illegible]

GDB your program

- GNU Project Debugger
- *b 5*
 - add a breakpoint at strcpy() in foo()
- *b 10*
 - add a breakpoint at char *str = "AAAA...";
- *r*
 - run to the breakpoint 2
- *info frame*

```
gdb-peda$ b 5
Breakpoint 1 at 0x8048411: file bufferoverflow.c, line 5.
gdb-peda$ b 10
Breakpoint 2 at 0x8048437: file bufferoverflow.c, line 10.
```

```
gdb-peda$ info frame
Stack level 0, frame at 0xbfffec90:
  eip = 0x8048437 in main (bufferoverflow.c:10); saved eip = 0xb7d82637
  called by frame at 0xbfffed00
  source language c.
  Arglist at 0xbfffec78, args:
  Locals at 0xbfffec78, Previous frame's sp is 0xbfffec90
  Saved registers:
  ebp at 0xbfffec78, eip at 0xbfffec8c
```


GDB your program

- *info registers*
 - check register values
- *x/16x 0xbfffece8*
 - See stack layout

```
gdb-peda$ info register
eax          0xb7f1ddbc      0xb7f1ddbc
ecx          0xbfffec90      0xbfffec90
edx          0xbfffecb4      0xbfffecb4
ebx          0x0            0x0
esp          0xbfffec60      0xbfffec60
ebp          0xbfffec78      0xbfffec78
esi          0xb7f1c000      0xb7f1c000
edi          0xb7f1c000      0xb7f1c000
eip          0x8048437        0x8048437 <main+17>
```

previous eip

```
gdb-peda$ x/16x 0xbfffec78
0xbfffec78: 0x00000000 0xb7d82637 0xb7f1c000 0xb7f1c000
0xbfffec88: 0x00000000 0xb7d82637 0x00000001 0xbfffed24
0xbfffec98: 0xbfffed2c 0x00000000 0x00000000 0x00000000
0xbfffece8: 0xb7f1c000 0xb7fffc04 0xb7fff000 0x00000000
```

- *C*
 - continues to run to next breakpoint (breakpoint 1)

- *info frame*

```
gdb-peda$ info frame
Stack level 0, frame at 0xbfffec50:
  eip = 0x8048411 in foo (bufferoverflow.c:5); saved eip = 0x8048449
  called by frame at 0xbfffec90
  source language c.
  Arglist at 0xbfffec48, args: str=0x80484e0 'A' <repeats 24 times>
  Locals at 0xbfffec48, Previous frame's sp is 0xbfffec50
  Saved registers:
    ebp at 0xbfffec48, eip at 0xbfffec4c
```


GDB your program

- *x/16x 0xbfffec48*

```
gdb-peda$ x/16x 0xbfffec48
0xbfffec48: 0xbfffec78 0x08048449 0x080484e0 0x00000000
0xbfffec58: 0xb7d98a50 0x080484ab 0x00000001 0xbfffed24
0xbfffec68: 0xbfffed2c 0x080484e0 0xb7f1c3dc 0xbfffec90
0xbfffec78: 0x00000000 0xb7d82637 0xb7f1c000 0xb7f1c000
```

previous eip

- *b 6*

- *c*

- *x/16x 0xbfffec48*

```
gdb-peda$ x/16x 0xbfffec48
0xbfffec48: 0x41414141 0x08048400 0x080484e0 0x00000000
0xbfffec58: 0xb7d98a50 0x080484ab 0x00000001 0xbfffed24
0xbfffec68: 0xbfffed2c 0x080484e0 0xb7f1c3dc 0xbfffec90
0xbfffec78: 0x00000000 0xb7d82637 0xb7f1c000 0xb7f1c000
```

previous EBP has
been mangled!

```
-----code-----
0x804841a <foo+15>: push    eax
0x804841b <foo+16>: call   0x80482e0 <strcpy@plt>
0x8048420 <foo+21>: add    esp,0x10
=> 0x8048423 <foo+24>: nop
0x8048424 <foo+25>: leave
0x8048425 <foo+26>: ret
0x8048426 <main>: lea    ecx,[esp+0x4]
0x804842a <main+4>: and    esp,0xffffffff
```

we have executed
strcpy()

GDB your program

- We have successfully mangled saved old EBP
- How to mangle return address?
 - return to the function you want to execute
 - (by carefully calculation and overflow)
- Solve it in our lab 1. (will release soon)

How to defend buffer overflow

- Lack checking? Do checking!
 - strcpy → strcpy_s in MS
 - <https://docs.microsoft.com/en-us/cpp/c-runtime-library/reference/strcpy-s-wcscpy-s-mbscpy-s?view=vs-2017>
- Address Randomization - ASLR
 - Attacker needs know where the stack is.
 - Make “guess” much harder
- StackGuard
 - Canary

Appendix



- Stack Guard Protection Scheme
 - Turned on by default
 - To turn it off: `gcc -fno-stack-protector example.c`
- Stack Overflow

... local variable ...	old ebp	return addr
...AAAAAA...	AAAA	AAAA

- Use “Canaries” to detect
 - A canary is a value inserted between *local variable* and *return address*
 - If you want to overwrite *return address*, the canary will be first overwritten
 - Detected !

-fno-stack-protector

- `gcc -fno-stack-protector -S -masm=intel -m32 -ggdb sample.c -o sample2.S`

```
add:
.LFB0:
.file 1 "sample.c"
.loc 1 2 0
.cfi_startproc
push    ebp
.cfi_def_cfa_offset 8
.cfi_offset 5, -8
mov     ebp, esp
.cfi_def_cfa_register 5
sub     esp, 2064
.loc 1 4 0
mov     edx, DWORD PTR [ebp+8]
mov     eax, DWORD PTR [ebp+12]
add     eax, edx
mov     DWORD PTR [ebp-4], eax
.loc 1 5 0
mov     eax, DWORD PTR [ebp-4]
.loc 1 6 0
leave
.cfi_restore 5
.cfi_def_cfa 4, 4
ret
.cfi_endproc
```

$16B + 512 * 4 = 2064B$, not know why 16B?
→ Stack Alignment (default 16B)

Parameter a

Parameter b

-fno-stack-protector

- `gcc -S -masm=intel -m32 -ggdb sample.c -o sample2.S`

```
add:
.LFB0:
.file 1 "sample.c"
.loc 1 2 0
.cfi_startproc
push    ebp
.cfi_def_cfa_offset 8
.cfi_offset 5, -8
mov     ebp, esp
.cfi_def_cfa_register 5
sub     esp, 2072
.loc 1 2 0
mov     eax, DWORD PTR gs:20
mov     DWORD PTR [ebp-12], eax
xor     eax, eax
.loc 1 4 0
mov     edx, DWORD PTR [ebp+8]
mov     eax, DWORD PTR [ebp+12]
add     eax, edx
mov     DWORD PTR [ebp-2064], eax
.loc 1 5 0
mov     eax, DWORD PTR [ebp-2064]
.loc 1 6 0
mov     ecx, DWORD PTR [ebp-12]
xor     ecx, DWORD PTR gs:20
je      .L3
call    __stack_chk_fail
```

$16B + 512 * 4 = 2064B$, why 2072B?!
8B for Canary !



canary is from `ebp-12` to `ebp-4`
DWORD = 8bytes

local variable `c` comes here
from `ebp-2064` to `ebp-2060` (4B)

buffer `temp[512]` is from
`ebp-2060` to `ebp-12` (2048B)

XOR to check canary